

Hydro Model Architectures


autopath

1 Overview

The **Hydro** datablock provides a flexible framework for latent-to-image decoders. It is designed to invert feature extractors (e.g., GigaPath) by mapping their embeddings back to the original image space. The system supports multiple architectural versions, dispatched via the **version** parameter in the model configuration.

2 Model_0: Deconvolutional Decoder

Model_0 is the baseline architecture, using progressive upsampling to reconstruct $N \times N$ RGB images from K -dimensional latent vectors.




2.1 Architectural Improvements to Combat Blur

To address the blurriness observed in **Model_0** outputs, several optional components can be enabled. These features prioritize the preservation and enhancement of high-frequency information.

2.1.1 Spatial Attention Gates


The **SpatialAttentionGate** identifies and emphasizes salient spatial features (e.g., edges and textures) while suppressing redundant background information.



By applying a learned spatial mask M , the model focuses its representational capacity on sharp edge transitions, effectively reducing the “halos” and washed-out textures typical of standard deconvolutional decoders.

2.1.2 Feature Modulation


Feature modulation allows the global latent vector $\mathbf{z}$ to dynamically scale the intensity of intermediate feature maps. This provides a mechanism for the latent representation to globally influence spatial details at every upsampling stage.



This global conditioning ensures that the reconstructed image maintains sharpening consistency across different scales, as the latent vector can “boost” specific frequency components needed for a sharp reconstruction.

2.1.3 Residual Upsampling

Residual upsampling incorporates long-range skip connections around the upsampling layers. This allows the network to learn only the residual difference needed to reach the next resolution, facilitating better gradient flow and preserving high-frequency details from lower spatial resolutions.




By bypassing the non-linear transformations for a portion of the signal, the model avoids the smoothing effect of repeated convolutions, resulting in crisper results.

3 Model_1: Vision Transformer (ViT) Decoder

Model_1 adopts a transformer-based architecture to model long-range dependencies and produce sharper reconstructions by operating on a tokenized representation of the image.

3.1 Architecture Overview

The architecture follows a “latent token to patch” approach:



3.1.1 Tokenization and Projection

The latent vector $\mathbf{z}$ is projected into a sequence of $M = (N/p)^2$ tokens, where p is the patch size (e.g., 16).

$$\mathbf{z} \xrightarrow{\text{MLP}} \mathbf{T} \in \mathbb{R}^{M \times D}$$

Learnable 2D positional embeddings are added to provide spatial context to the transformer.

3.1.2 Transformer Backbone

A series of L standard Transformer blocks (Multi-Head Self-Attention + MLP) process the token sequence. This allows information to propagate globally across the image grid.

3.1.3 Reconstruction Head

Each token $\mathbf{t}_i \in \mathbb{R}^D$ is projected to a patch $\mathbf{P}_i \in \mathbb{R}^{3 \times p \times p}$ using a linear layer. The patches are then rearranged to form the final image $\mathbf{I} \in \mathbb{R}^{3 \times N \times N}$.

4 System Configuration

The Hydro datablock uses a shared configuration dataclass:

Parameter	Default	Description
latent_dim	1536	Input latent vector dimension K
image_size	256	Output resolution $N \times N$
hidden_channels	256	Base channel count (Model_0) / Dimension D (Model_1)
patch_size	16	Patch size for Model_1
num_layers	12	Number of layers (Transformer blocks) for Model_1
version	0	Model architecture version (0 or 1)
use_residual_upsampling	False	Enable residual skip connections in Model_0
use_attention_gates	False	Enable spatial attention in Model_0
use_feature_modulation	False	Enable latent-based feature modulation in Model_0
loss_type	mse	MSE, L1, L0, or LPIPS reconstruction loss

5 Loss Functions

Reconstruction quality is improved by incorporating perceptual metrics:

$$\mathcal{L} = \begin{cases} \frac{1}{3N^2} \|\hat{\mathbf{I}} - \mathbf{I}^*\|_2^2 & \text{if loss_type = "mse"} \\ \frac{1}{3N^2} \|\hat{\mathbf{I}} - \mathbf{I}^*\|_1 & \text{if loss_type = "l1"} \\ \text{LPIPS}(\hat{\mathbf{I}}, \mathbf{I}^*) & \text{if loss_type = "lpips"} \end{cases}$$

The **LPIPS** (Learned Perceptual Image Patch Similarity) loss compares deep features from a pre-trained network (e.g., VGG) and is highly effective at reducing blur by penalizing differences in perceptual content rather than raw pixel values.